

How to use Mixture in your projects

Mixture is a deliberately-small, layered skills + agents framework for **Claude Code**. This guide shows how to drop it into another project and use each layer. Nothing here is theoretical — every command below is in this repo and tested.

Prerequisites: *Node.js ≥ 18, Claude Code. No npm dependencies — the tooling is pure Node built-ins.*

TL;DR — install with the CLI

From inside the project you want to add Mixture to:

```
# from npm (published):
npx mixture-skills install --profile dev

# or pin straight from source:
npx github:dikshant-goforsys/mixture install --profile dev
```

Common variants:

```
npx mixture list                # show profiles and their skills
npx mixture install --profile full --with-memory --with-coordination
npx mixture doctor              # check what's installed here
npx mixture install --global --link # symlink into ~/.claude/skills
(dev on a clone)
```

This installs skills into `<project>/.claude/skills/` (or `~/.claude/skills` with `--global`), and with the flags also drops the memory/coordination runtimes under `<project>/.mixture/framework/` and wires the memory hooks into `.claude/settings.json` (backed up first; re-running is idempotent). Then restart Claude Code or `/reload-skills`.

Manual alternative (no CLI): `cp -r /path/to/mixture/skills/<bucket>/<skill> your-project/.claude/skills/`.

That's enough to get the **skills** (L1–L2). L3 (hooks/memory) and L4 (coordination) are opt-in flags — details below.

The four layers (what you're getting)

Layer	What it is	Where
L1 Behavioral kernel	One always-relevant coding-discipline skill	skills/kernel/coding-behavior
L2 Skills + authoring gate	grill-me, diagnose, tdd, context, code-review, dev-loop, frontend-design + write-a-skill	skills/engineering/*, skills/meta/*
L3 Governance	Install profiles, env-governed hooks, session memory, CI gate	manifests/, hooks/, scripts/

L4 Coordination	Multi-agent task ledger + heartbeat runtime	coordination/, skills/coordination/*
-----------------	---	---

Lower layers never depend on higher ones. Use only the layers you want.

Installing the skills (L1–L2)

A "skill" is a directory with a `SKILL.md`. Claude Code auto-loads a skill when your request matches its `description`. Three ways to install:

A. Global symlink (all your projects get them): the TL;DR loop above, linking into `~/.claude/skills/`.

B. Per-project: copy the skill dirs you want into `your-project/.claude/skills/`:

```
cp -r /path/to/mixture/skills/engineering/diagnose your-project/.claude/skills/
```

C. As a plugin: the repo ships `.claude-plugin/plugin.json` listing all skills, so it can be installed as a Claude Code plugin (`/plugin`) if you host it in a marketplace.

Pick a subset with profiles (`manifests/install-profiles.json`): `minimal` (kernel only), `dev` (kernel

- `engineering`), `frontend` (dev loop + shadcn/ui design system), `authoring`, `coordination`, `full`. Install only the skills a profile lists to keep your context budget small.

Using the skills

Just work normally — Claude routes to a skill by its `description`. Or invoke explicitly by name:

- **`coding-behavior`** (kernel) — fires on real coding tasks; biases toward asking before assuming, minimal diffs, surgical changes, and test-verifiable goals.
- **`grill-me`** — before a big build, interrogates requirements one question at a time.
- **`diagnose`** — for bugs: builds a deterministic repro *first*, then tests ranked hypotheses.
- **`tdd`** — vertical-slice red-green-refactor.
- **`context`** — maintains a `CONTEXT.md` domain glossary (shared human↔agent vocabulary).
- **`code-review`** — prioritized correctness/security findings with `file:line`, not style nits.
- **`dev-loop`** — the delivery contract: `tdd` → full gate (tests/lint/typecheck) → code-review verdict, looped until the greenzone (everything green + zero must-fix findings).
- **`frontend-design`** — distinctive, production-grade UI (no generic "AI slop" aesthetics) executed on the shadcn/ui system: CLI-added owned components, Radix accessibility, semantic tokens.

Governance (L3) — optional but recommended

Install profiles

`manifests/install-profiles.json` maps a profile name → a set of skill paths. Use it to decide what to install per project (a heavy repo can run `minimal`).

Hooks, governed by env (never by editing files)

`hooks/hooks.json` is the annotated source of truth. Resolve it into a Claude Code-consumable config:

```
node scripts/resolve-hooks.mjs          # prints the active hook config for your
profile
node scripts/resolve-hooks.mjs --explain  # shows which hooks are on/off and why
```

Control behavior with env vars — no file edits:

- `MIXTURE_HOOK_PROFILE=off|standard|strict` (default `standard`)
- `MIXTURE_DISABLED_HOOKS=load-memory,save-memory` (comma-separated hook names)

To wire into Claude Code, put the resolved output's `hooks` into your project `.claude/settings.json`.

*Note: `hooks/hooks.json` commands are repo-relative — this flow is for working **on the Mixture repo itself**. In a consumer project, don't resolve this file; `npx mixture-skills install --with-memory` wires the memory hooks with consumer-correct paths (`.mixture/framework/memory/*`) automatically.*

Session memory (survives compaction)

Durable facts live outside the context window. **Two backends**, selected by `MIXTURE_MEMORY_BACKEND` :

- `json` (default) — a single `.mixture/memory/store.json` ; works on Node ≥ 18 .
- `sqlite` — a local `.mixture/memory/store.db` via the **built-in `node:sqlite`** (zero extra deps; needs Node ≥ 22).

Install with SQLite wired into the hooks:

```
npx github:dikshant-goforsys/mixture install --profile dev --with-memory --memory-
backend sqlite
```

The CLI is backend-agnostic — same commands either way:

```
# add a durable fact (or wire as a tool the agent calls)
node hooks/memory-persistence/save.mjs --add "Chose Postgres over SQLite for multi-
writer support" --type decision --pin
node hooks/memory-persistence/load.mjs      # SessionStart: prints a digest for
context injection
node hooks/memory-persistence/save.mjs      # PreCompact: consolidate + dedupe
(idempotent)
node hooks/memory-persistence/clean.mjs     # SessionEnd: enforce TTL + max (never
evicts pinned)
```

Types: `goal` | `decision` | `constraint` | `glossary` | `fact` . Tunables:

`MIXTURE_MEMORY_{DIR,LOAD_LIMIT,MAX,TTL_DAYS}` .

The CI gate (keep the catalog honest)

```
npm run ci  # = validate-frontmatter + check-drift --ci + resolve-hooks --check +
coordination tests
```

- `validate-frontmatter.mjs` — every `SKILL.md` has a valid routing contract (kebab name, ≤ 1024 -char description containing "Use when").

- `check-drift.mjs` — every shipped skill is registered + in a profile + has an eval, and the catalog is under the ~30 cap. **In-editor (advisory) it warns; with `--ci` it's fatal.** (See ADR-0002.)

Copy `.github/workflows/ci.yml` to gate this on every PR.

Coordination (L4) — multi-agent runtime, optional

Use this only if you actually run multiple agents that share work. It's a task ledger with every safety invariant enforced in code.

Task ledger CLI

```
export MIXTURE_COORD_DIR=.mixture/coordination # where the ledger lives (gitignore it)

node coordination/cli.mjs create --title "Design schema" --priority 1
node coordination/cli.mjs create --title "Build API" --priority 1
node coordination/cli.mjs block --id T-2 --by T-1 # T-2 waits on T-1
node coordination/cli.mjs tick --agent worker-1 --run r1 # atomically claim top-priority ready work
node coordination/cli.mjs done --id T-1 # auto-resumes T-2
node coordination/cli.mjs ask --id T-2 --kind request_confirmation --prompt "Ship?" --key ship:T2:r1
node coordination/cli.mjs resolve --iid I-1 --response yes # wakes the assignee
```

Exit codes (machine-detectable — the agent never parses prose)

0 ok · 2 usage · 3 lock-busy · 5 stale-gate · 6 cycle · 7 budget-exhausted · 8 blocked · 9

CONFLICT → another agent owns it; **STOP**, never retry.

The behavioral contract

Install the `coordination-protocol` skill — it tells a woken agent exactly how to behave (identity → pick work → atomic checkout → do → delegate → human-gate), and to obey the exit codes. Hard invariants live in `coordination/ledger.mjs`; the skill is the behavior.

Run a heartbeat (Claude Code native)

- **Scheduled:** create a cron routine that fires every N minutes and runs `cli.mjs tick`, then executes any assigned task. (We use a recurring job at `:17,:47`.)
- **Self-paced:** an agent loops via `ScheduleWakeup` after each tick.
- **Dispatch many:** fan out one subagent per ready task with the `Agent` tool — atomic checkout guarantees no two touch the same task. (Proven in `coordination/DEMO.md`: a 4-agent contention test produced exactly one winner, no double-work.)

Extending: add your own skill (the right way)

Always go through the gate skill `write-a-skill`. The rules it enforces:

1. **description** = **routing contract**: ≤1024 chars, sentence 1 = what it does, sentence 2 = "Use when ...".

2. **Progressive disclosure:** `SKILL.md` ≤ ~100 lines; deeper detail in `references/*.md` one level deep.
3. **Earn your context** — prevents a specific failure the kernel/existing skills don't. Catalog cap ~30.
4. **Examples-as-docs:** at least one /  pair.
5. **Mandatory eval:** add `evals/<skill-name>.eval.md`.
6. **Register:** add to `.claude-plugin/plugin.json` and ≥1 profile.

Then: `node scripts/validate-frontmatter.mjs && node scripts/check-drift.mjs --ci` must pass.

Authoring tip: write the `SKILL.md` first, then register, then add the eval — the in-editor drift hook is *advisory* so it won't block this order; CI enforces completeness at merge (ADR-0002).

Environment variables reference

Var	Default	Purpose
MIXTURE_HOOK_PROFILE	standard	off/standard/strict — which hooks are active
MIXTURE_DISABLED_HOOKS	(none)	comma-separated hook names to disable
MIXTURE_MEMORY_DIR	.mixture/memory	memory store location
MIXTURE_MEMORY_BACKEND	json	json or sqlite (sqlite needs Node ≥ 22, zero extra deps)
MIXTURE_MEMORY_LOAD_LIMIT	50	max entries injected at SessionStart
MIXTURE_MEMORY_MAX / MIXTURE_MEMORY_TTL_DAYS	200 / 90	memory bound
MIXTURE_COORD_DIR	.mixture/coordination	task ledger location
MIXTURE_DRIFT_STRICT=1	—	make check-drift fatal (same as --ci)

Publishing to npm (for the maintainer)

So that anyone can run `npmx mixture-skills install ...`. You publish from the framework repo; consumers never do.

The bare name `mixture` is already taken on npm, so this package is `mixture-skills` (the CLI still answers to both `mixture-skills` and `mixture` once installed). `package.json` is already set up: `non-private`, `bin` wired, `files` whitelist, and a `prepublishOnly` that runs `npm run ci` so a failing gate blocks the publish.

One-time setup

1. Create a free npm account at npmjs.com (enable 2FA — recommended).
2. Log in on this machine (interactive — run it yourself):

```
npm login
```

Confirm with `npm whoami` .

Pre-flight (every release)

```
npm run ci          # gates must pass (prepublishOnly runs this again anyway)
npm pack --dry-run  # inspect exactly what will ship – should be ~30 files, no
.mixture/ state
```

Publish

```
npm version patch    # 0.1.0 -> 0.1.1 (use minor/major as appropriate); commits +
tags
npm publish          # unscoped + public by default
```

Done. From any project, anyone can now:

```
npx mixture-skills install --profile dev
```

Releasing updates

Bump and publish again — consumers get the new version on their next `npx` (or `npx mixture-skills@latest`):

```
npm version patch && npm publish
npx mixture-skills doctor # consumers can re-run install to refresh
```

If you prefer a scoped package

Use `@<your-npm-username>/mixture` instead (scoped names are always free, namespaced to you):

- set `"name": "@you/mixture"` in `package.json` , then `npm publish --access public`
- consumers run `npx @you/mixture install ...`

No npm at all (works today, zero setup)

The repo is on GitHub, so you can skip npm publishing entirely:

```
npx github:dikshant-goforsys/mixture install --profile dev
```

Keep the philosophy (or it rots)

Mixture's value is its discipline, not its size. If you adopt it, keep the guardrails:

- **The ~30-skill cap.** Adding one may mean deprecating one. Volume is the failure mode (see `docs/premortem.md`).
- **Every skill has an eval.** No eval = a guess; CI rejects it.

- **Enforce in code, not prose.** Invariants belong in `scripts/ / ledger.mjs` , not just instructions.
- **Don't build L4 before you need it.** It's powerful and risky; use the manual ledger until a real fleet exists.

See `docs/architecture.md` (what each idea was taken from), `docs/premortem.md` (10 ways it fails + guards), and `docs/adr/` (decisions) for the full reasoning.